

Signal-to-Pipeline Engine

How I built a full outbound system from scratch in 3 weeks and hit 1.5%+ reply rates at scale

CASE STUDY · RASUL SHAIKH · GTM ENGINEER

2 Wks → 3 Days

DEPLOY TIME

1.5%+

REPLY RATE

48h → 4h

QUALIFICATION CYCLE

6+ hrs → 30 min

RESEARCH/BATCH

THE PROBLEM

Omnibound needed to go from zero to a running outbound system — fast. No CRM data, no sequences, no enrichment pipeline. Manual research was taking 6+ hours per batch of 50 leads. Qualification required a senior rep to spend 48 hours per account. The team had no repeatable motion and no tooling to scale beyond 2 people.

THE ARCHITECTURE

- Signal layer: 6 data sources unified — LinkedIn job postings, funding data, tech stack signals, news events, web traffic spikes, and intent scores — merged into a single enriched record per account.
- AI scoring: Claude API rubric scoring each account 0–100 on firmographic fit, technographic alignment, and behavioral signals. Accounts below 60 are automatically skipped.
- Clay pipeline: 7-stage LLM waterfall for hyper-personalization — company summary, pain hypothesis, relevant case study match, personalized subject line, email body v1, email body v2 variant, and objection pre-emption.
- Enrichment agent: Python script (Playwright + asyncio) scraping 100+ domains headlessly, scoring accounts via the Claude API rubric, and auto-pushing qualified accounts to HubSpot with full routing context.
- Sending infrastructure: Smartlead across 250+ sending domains with inbox rotation, warmup schedules, and reply detection feeding back into the CRM sequence stage.
- Qualification layer: n8n workflow parsing replies using Claude, routing positive replies to a Calendly link and flagging referrals for manual follow-up within 4 hours.

Signal-to-Pipeline Engine

CASE STUDY · RASUL SHAIKH · GTM ENGINEER

THE RESULTS

- Reply rate: 1.5%+ sustained across 3 months and 12,000+ contacts — 3x the industry average for cold email.
- Deploy time: Full system (enrichment → scoring → Clay pipeline → Smartlead → CRM routing) from kick-off to first send in 3 weeks. Subsequent campaigns deploy in 3 days.
- Qualification speed: AI reply parsing + n8n routing cut manual qualification from 48 hours to 4 hours per positive reply.
- Research efficiency: Python enrichment agent replaced 6+ hours/batch of manual research with a sub-30-minute automated pass, running unattended.
- Pipeline generated: \$3.5M+ in qualified pipeline from the system over the engagement period.

KEY LEARNINGS

- Signal quality beats volume. A list of 500 accounts scored above 70 outperforms 5,000 unsegmented contacts every time.
- The personalization hook is the leverage point — not the email body. Getting the first line right (via Clay LLM + Playwright-scraped context) is what drives reply rates above 1%.
- Reply parsing is the bottleneck nobody talks about. Automating intent classification was what unlocked 4-hour qualification cycles.
- Infrastructure is GTM strategy. A rep working this system outperforms 3 reps without it.

STACK

Enrichment: Python (Playwright, asyncio), Supabase, Firecrawl, Prospeo

AI / LLM: Claude API (scoring + copy), Claude Code (agent orchestration), OpenAI API

Outbound: Clay (7-stage pipeline), Smartlead (250+ domains), HeyReach (LinkedIn)

Automation: n8n (reply routing), HubSpot (CRM + routing), Webhooks